



Gemini 2.5 Flash

Best for balancing reasoning and speed

Open in Agent Studio

View code

- Latest
- Overview
- Documentation



Provisioned throughput is available for Gemini 2.5 Flash

Assure your generative AI workloads for supported models on Agent Platform.

[Learn more](#)

Manage orders

Place order

Latest

Gemini 2.5 Flash is a thinking model, designed to tackle increasingly complex problems.

Key features:

- User can activate thinking
- User can set a thinking budget (8k-24k tokens)

Gemini 2.5 Flash is made available through our new [Gen AI SDK](#) that provides a unified interface for Google AI Studio and Agent Platform in Python and Go.

For more information about differences with the previous model version, see [Model versions and lifecycle](#).

Overview

Gemini 2.5 models are thinking models, capable of reasoning through their thoughts before responding, resulting in enhanced performance and improved accuracy.

Gemini 2.5 Flash: best for balancing reasoning and speed.

Model details

The model details are as follows:

Property	Description
Model name	gemini-2.5-flash
Supported data types	Inputs: Audio, images, video, PDF, and text Output: Text only
Token limits	Input token limit: 1,048,576 Output token limit: 64k